

```
In [1]: import pandas as pd
import seaborn as sns
import neattext.functions as nfx
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
import joblib
```

```
In [2]: df_train = pd.read_csv('train_150k.txt', sep='\t', names=['sent', 'tweet'])
df_test = pd.read_csv('test_62k.txt', sep='\t', names=['sent', 'tweet'])

print("Training Dataset\n", df_train.head())
print("\nTesting Dataset\n", df_test.head())
print("\nSent total count\n", df_train['sent'].value_counts())
```

Training Dataset

	sent	tweet
0	0	Starting back at work today Looks like it'l...
1	1	Sugar levels dropping... munchies setting in. ...
2	1	@karineb22 yeah!!! have a great summer break!
3	1	hannah montana was very good. now going to re...
4	1	@Mayra326 aww, have fun! I just had my 3D las...

Testing Dataset

	sent	tweet
0	1	@justineville ...yeahhh.) i'm 39 tweets from ...
1	0	@ApplesnFeathers aww. Poor baby! On your only ...
2	0	@joeymcintyre With my refunded \$225 (Australia...
3	0	It's fine. Today sucks just because me those t...
4	0	Im just chilling on psp and stuff, but sitting...

Sent total count

0	75019
1	74966

Name: sent, dtype: int64

```
In [3]: # Remove user handles
df_train['Clean_Tweet'] = df_train['tweet'].apply(nfx.remove_userhandles)
df_test['Clean_Tweet'] = df_test['tweet'].apply(nfx.remove_userhandles)

# Remove stopwords
df_train['Clean_Tweet'] = df_train['Clean_Tweet'].apply(nfx.remove_stopwords)
df_test['Clean_Tweet'] = df_test['Clean_Tweet'].apply(nfx.remove_stopwords)

print(df_train.head())
```

	sent	tweet \	Clean_Tweet
0	0	Starting back at work today Looks like it'll...	Starting work today Looks like it'll raining c...
1	1	Sugar levels dropping... munchies setting in. ...	Sugar levels dropping... munchies setting in. ...
2	1	@karineb22 yeah!!! have a great summer break!	yeah!!! great summer break!
3	1	hannah montana was very good. now going to re...	hannah montana good. going read twilight.
4	1	@Mayra326 aww, have fun! I just had my 3D las...	aww, fun! 3D week!

```
In [4]: X_train = df_train['Clean_Tweet']
y_train = df_train['sent']

X_test = df_test['Clean_Tweet']
y_test = df_test['sent']
```

```
In [5]: X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.2, random_state=42)
```

```
In [6]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import Perceptron
from sklearn.metrics import accuracy_score

# Convert text to numerical features using TF-IDF
vectorizer = TfidfVectorizer(max_features=5000)
X_train_tfidf = vectorizer.fit_transform(X_train)
X_val_tfidf = vectorizer.transform(X_val) # Use validation set

# Define and train the Perceptron model
perceptron = Perceptron(max_iter=1000, tol=1e-3, random_state=42)
perceptron.fit(X_train_tfidf, y_train)

# Make predictions
y_pred = perceptron.predict(X_val_tfidf)

# Evaluate the model
accuracy = accuracy_score(y_val, y_pred)
print("Perceptron Accuracy:", accuracy)

Perceptron Accuracy: 0.6903058049604409
```

```
In [7]: from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression

# Updated: Use TF-IDF instead of CountVectorizer
pipe_lr = Pipeline(steps=[
    ('tfidf', TfidfVectorizer(max_features=50000)),
    ('lr', LogisticRegression(max_iter=200))
])

# Train the model
pipe_lr.fit(X_train, y_train)

# Evaluate the model
print("Logistic Regression Accuracy:", pipe_lr.score(X_val, y_val))
```

Logistic Regression Accuracy: 0.7661125433371855

```
In [8]: pipe_lr = Pipeline(steps=[('cv', CountVectorizer()), ('lr', LogisticRegression())])
pipe_lr.fit(X_train, y_train)
print("Logistic Regression Accuracy:", pipe_lr.score(X_val, y_val))
```

/Users/sakshi/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(
Logistic Regression Accuracy: 0.7620010667614899

```
In [9]: from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.svm import LinearSVC

# Define pipeline
pipe_svm = Pipeline(steps=[
    ('tfidf', TfidfVectorizer()),
    ('svc', LinearSVC())
])

# Train the model
pipe_svm.fit(X_train, y_train)

# Evaluate the model
print("SVM Accuracy:", pipe_svm.score(X_val, y_val))
```

SVM Accuracy: 0.7508445195128456

```
In [10]: pipe_rf = Pipeline(steps=[('cv', CountVectorizer()), ('rf', RandomForestClassifier())])
pipe_rf.fit(X_train, y_train)
print("Random Forest Accuracy:", pipe_rf.score(X_val, y_val))
```

Random Forest Accuracy: 0.7250644501733488


```

In [11]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import Perceptron, LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve
from sklearn.preprocessing import label_binarize

# Reduce TF-IDF features for faster computation
vectorizer = TfidfVectorizer(max_features=5000)
X_train_tfidf = vectorizer.fit_transform(X_train[:10000]) # Training
X_val_tfidf = vectorizer.transform(X_val)

# Initialize models with optimized parameters
models = {
    "Perceptron": Perceptron(max_iter=500, tol=1e-3, random_state=42),
    "Logistic Regression": LogisticRegression(max_iter=500, random_state=42),
    "SVM": SVC(kernel='linear', C=1, probability=True), # probability
    "Random Forest": RandomForestClassifier(n_estimators=5, random_state=42)
}

# Train and evaluate models
accuracies = {}
conf_matrices = {}
roc_curves = {}
roc_aucs = {}

# Convert labels to binary format for ROC curve
y_val_bin = label_binarize(y_val, classes=[0, 1]) # Assuming binary classes

for name, model in models.items():
    model.fit(X_train_tfidf, y_train[:10000]) # Training on subset
    y_pred = model.predict(X_val_tfidf)

    # Store accuracy
    accuracies[name] = accuracy_score(y_val, y_pred)

    # Store confusion matrix
    conf_matrices[name] = confusion_matrix(y_val, y_pred)

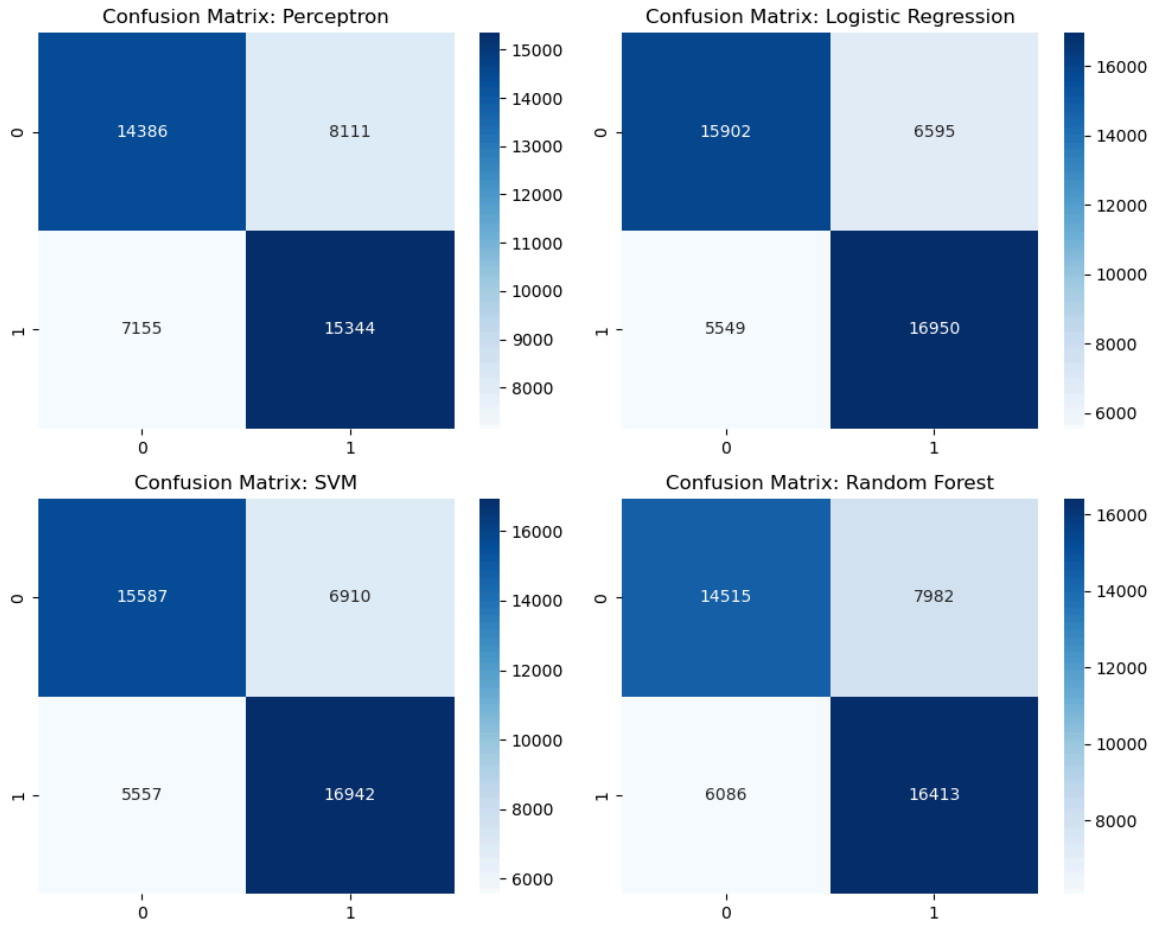
    # ROC Curve
    if hasattr(model, "predict_proba"): # Check if model supports probability
        y_proba = model.predict_proba(X_val_tfidf)[:, 1]
        fpr, tpr, _ = roc_curve(y_val, y_proba)
        roc_auc = auc(fpr, tpr)
        roc_curves[name] = (fpr, tpr)
        roc_aucs[name] = roc_auc

# Convert to DataFrame
acc_df = pd.DataFrame(accuracies, index=["Accuracy"]).T

# ===== 2 CONFUSION MATRICES =====
fig, axes = plt.subplots(2, 2, figsize=(10, 8))
axes = axes.flatten()
for i, (name, matrix) in enumerate(conf_matrices.items()):
    sns.heatmap(matrix, annot=True, fmt='d', cmap='Blues', ax=axes[i])
    axes[i].set_title(f"Confusion Matrix: {name}")
plt.tight_layout()

```

```
plt.show()
```




```

In [12]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import Perceptron, LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve
from sklearn.preprocessing import label_binarize

# Reduce TF-IDF features for faster computation
vectorizer = TfidfVectorizer(max_features=5000)
X_train_tfidf = vectorizer.fit_transform(X_train[:10000]) # Training
X_val_tfidf = vectorizer.transform(X_val)

# Initialize models
models = {
    "Perceptron": Perceptron(max_iter=500, tol=1e-3, random_state=42),
    "Logistic Regression": LogisticRegression(max_iter=500, random_state=42),
    "SVM": SVC(kernel='linear', C=1, probability=True),
    "Random Forest": RandomForestClassifier(n_estimators=5, random_state=42)
}

# Train and evaluate models
accuracies = {}
conf_matrices = {}
roc_curves = {}
roc_aucs = {}

y_val_bin = label_binarize(y_val, classes=[0, 1])

for name, model in models.items():
    model.fit(X_train_tfidf, y_train[:10000])
    y_pred = model.predict(X_val_tfidf)

    # Accuracy
    accuracies[name] = accuracy_score(y_val, y_pred)

    # Confusion Matrix
    conf_matrices[name] = confusion_matrix(y_val, y_pred)

    # ROC Curve
    if hasattr(model, "predict_proba"):
        y_proba = model.predict_proba(X_val_tfidf)[:, 1]
        fpr, tpr, _ = roc_curve(y_val, y_proba)
        roc_auc = auc(fpr, tpr)
        roc_curves[name] = (fpr, tpr)
        roc_aucs[name] = roc_auc

# Accuracy Table
acc_df = pd.DataFrame(accuracies, index=["Accuracy"]).T

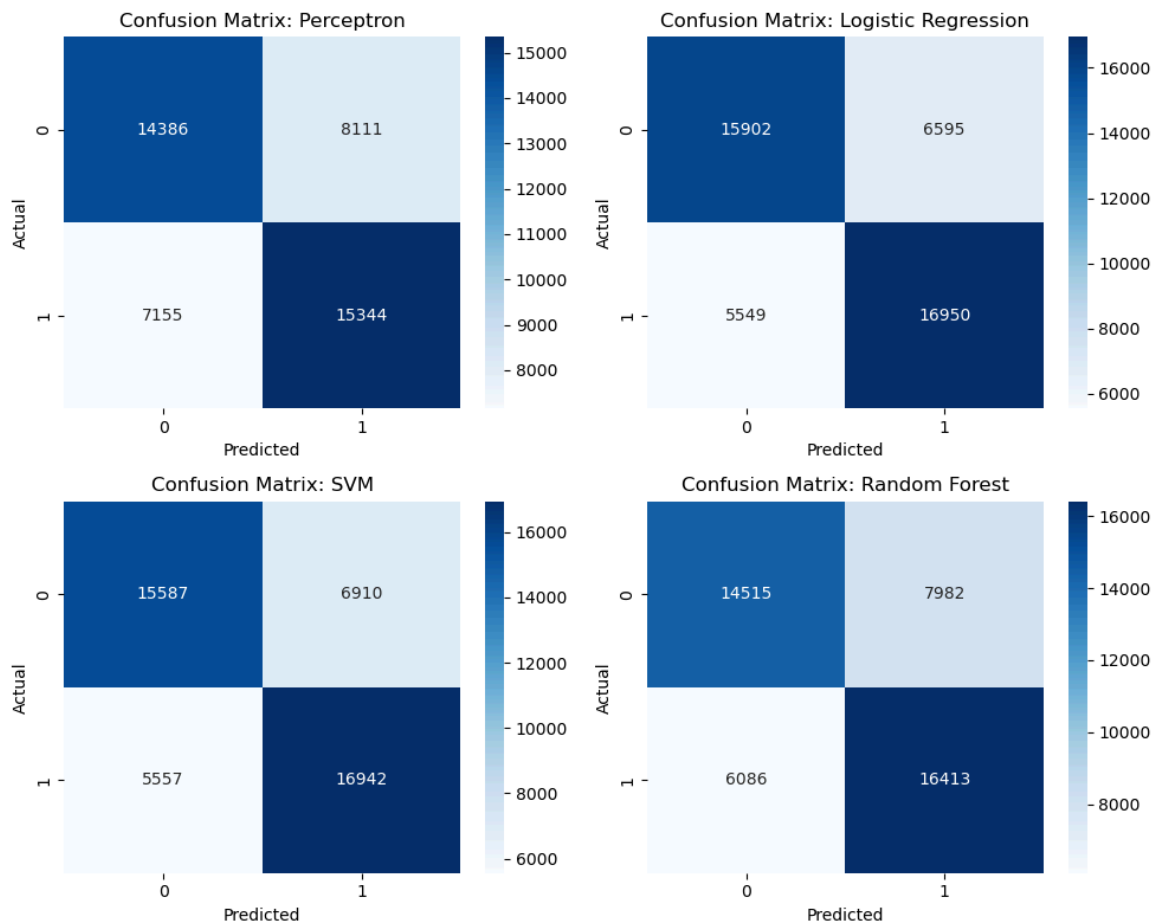
# 1 Plot Individual Confusion Matrices
fig, axes = plt.subplots(2, 2, figsize=(10, 8))
axes = axes.flatten()
for i, (name, matrix) in enumerate(conf_matrices.items()):
    sns.heatmap(matrix, annot=True, fmt='d', cmap='Blues', ax=axes[i])
    axes[i].set_title(f"Confusion Matrix: {name}")
    axes[i].set_xlabel("Predicted")
    axes[i].set_ylabel("Actual")

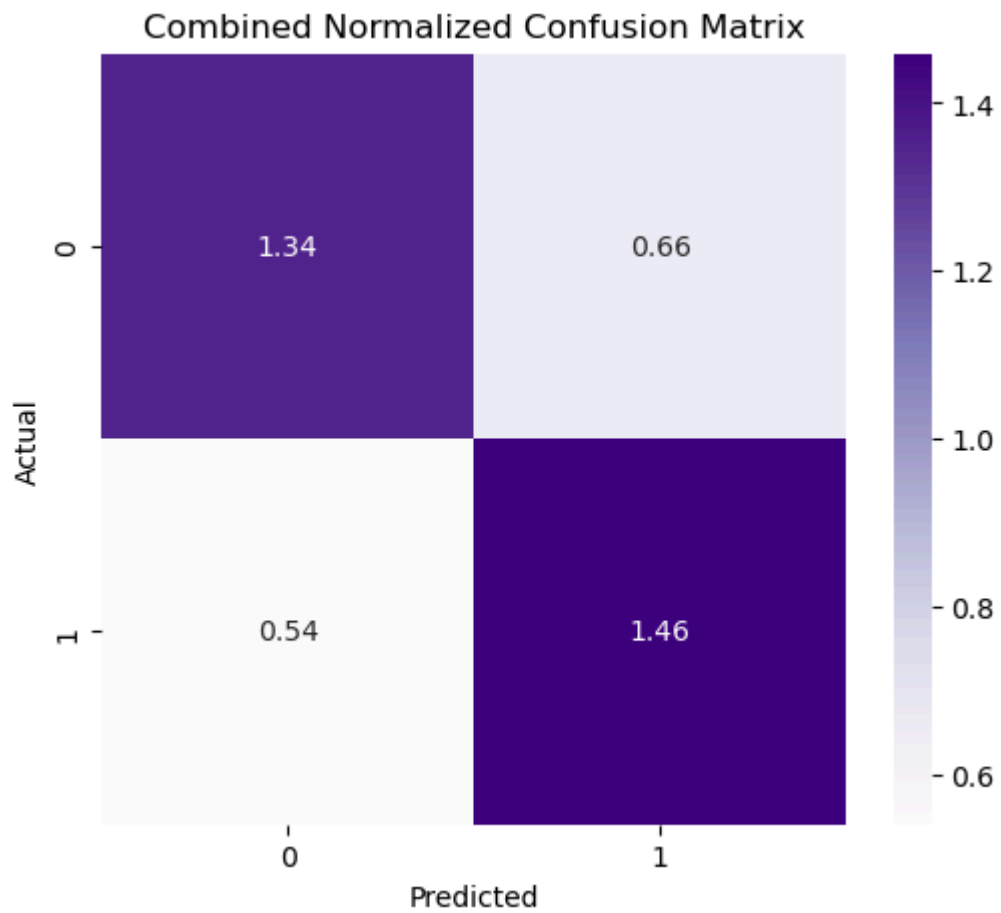
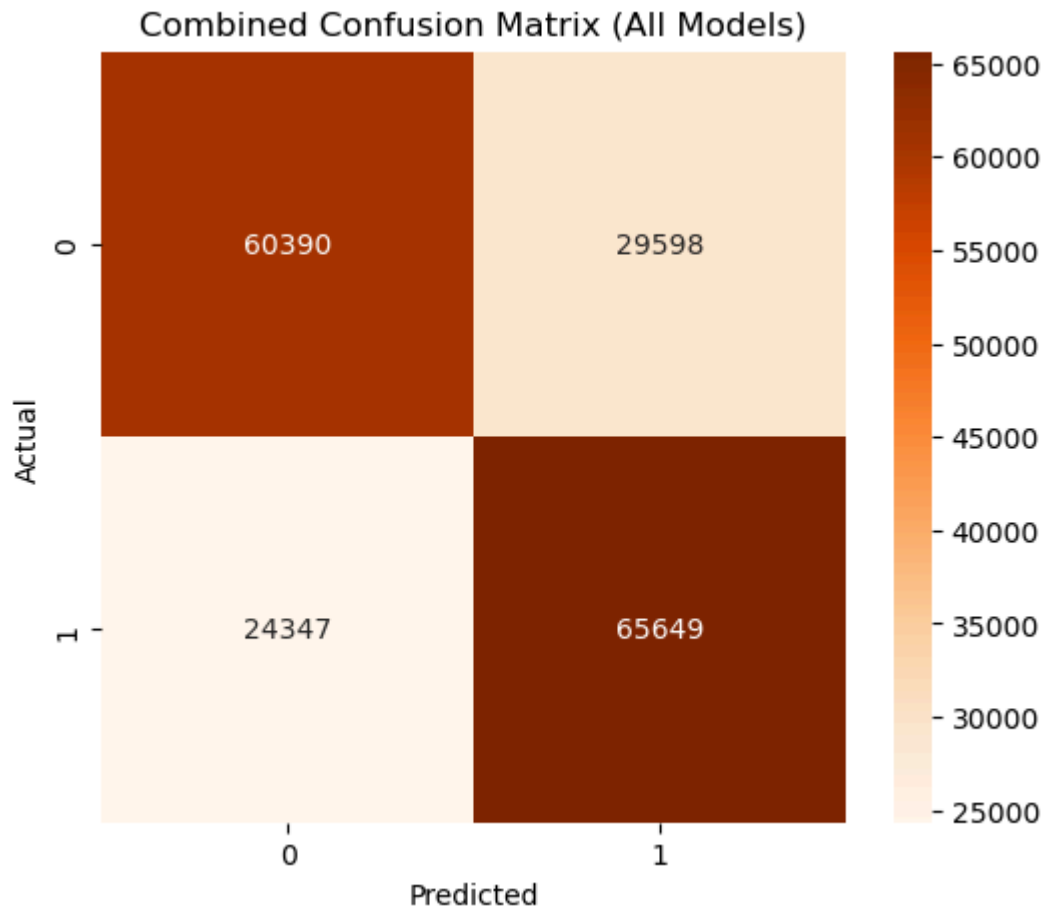
```

```
plt.tight_layout()
plt.show()

# 2 Combined Confusion Matrix (Raw Sum)
combined_conf_matrix = sum(conf_matrices.values())
plt.figure(figsize=(6, 5))
sns.heatmap(combined_conf_matrix, annot=True, fmt='d', cmap='Oranges')
plt.title("Combined Confusion Matrix (All Models)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# 3 Combined Normalized Confusion Matrix
normalized_matrices = [matrix / matrix.sum() for matrix in conf_matrices.values()]
combined_normalized = sum(normalized_matrices)
plt.figure(figsize=(6, 5))
sns.heatmap(combined_normalized, annot=True, fmt='.2f', cmap='Purples')
plt.title("Combined Normalized Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```






```

In [13]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import Perceptron, LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve
from sklearn.preprocessing import label_binarize

# Reduce TF-IDF features for faster computation
vectorizer = TfidfVectorizer(max_features=5000)
X_train_tfidf = vectorizer.fit_transform(X_train[:10000]) # Training
X_val_tfidf = vectorizer.transform(X_val)

# Initialize models with optimized parameters
models = {
    "Perceptron": Perceptron(max_iter=500, tol=1e-3, random_state=42),
    "Logistic Regression": LogisticRegression(max_iter=500, random_state=42),
    "SVM": SVC(kernel='linear', C=1, probability=True), # probability
    "Random Forest": RandomForestClassifier(n_estimators=5, random_state=42)
}

# Train and evaluate models
accuracies = {}
conf_matrices = {}
roc_curves = {}
roc_aucs = {}

# Convert labels to binary format for ROC curve
y_val_bin = label_binarize(y_val, classes=[0, 1]) # Assuming binary classes

for name, model in models.items():
    model.fit(X_train_tfidf, y_train[:10000]) # Training on subset
    y_pred = model.predict(X_val_tfidf)

    # Store accuracy
    accuracies[name] = accuracy_score(y_val, y_pred)

    # Store confusion matrix
    conf_matrices[name] = confusion_matrix(y_val, y_pred)

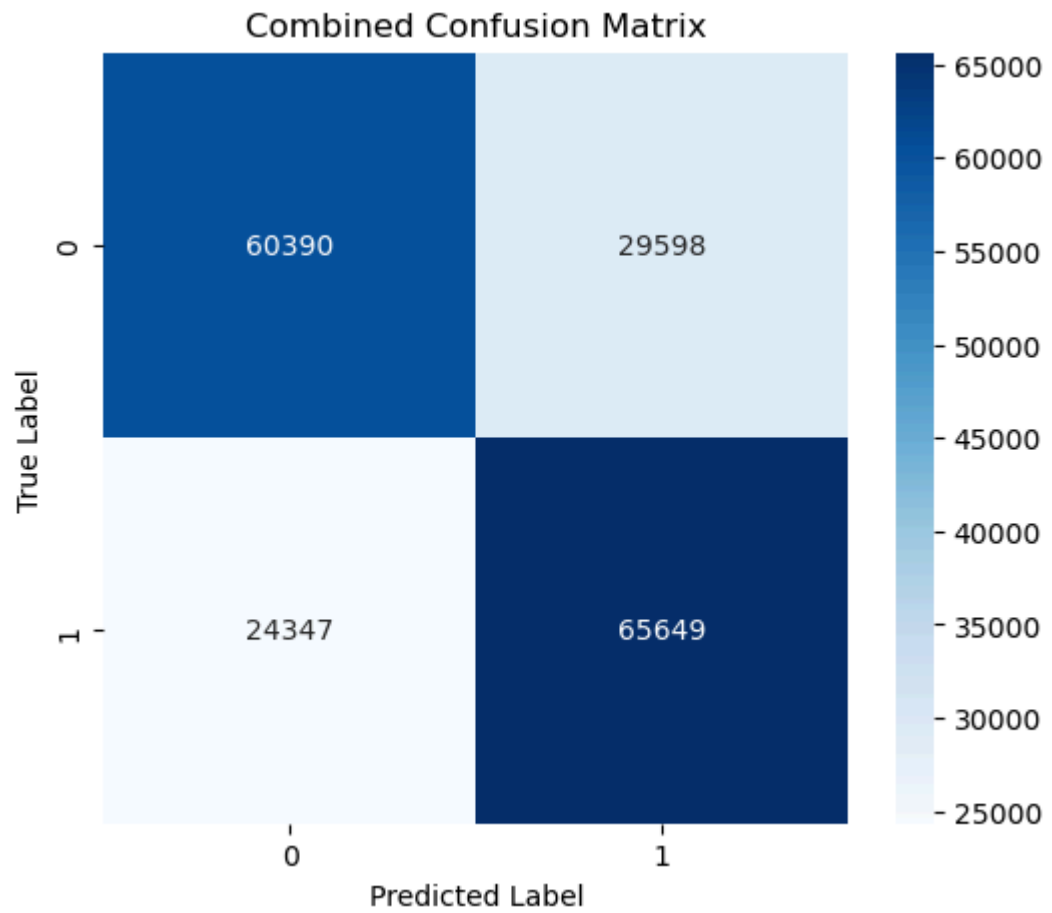
    # ROC Curve
    if hasattr(model, "predict_proba"): # Check if model supports probability
        y_proba = model.predict_proba(X_val_tfidf)[:, 1]
        fpr, tpr, _ = roc_curve(y_val, y_proba)
        roc_auc = auc(fpr, tpr)
        roc_curves[name] = (fpr, tpr)
        roc_aucs[name] = roc_auc

# Combine confusion matrices by summing them
combined_matrix = sum(conf_matrices.values())

# Plot combined confusion matrix
plt.figure(figsize=(6, 5))
sns.heatmap(combined_matrix, annot=True, fmt='d', cmap='Blues')
plt.title("Combined Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")

```

```
plt.show()
```



```

In [14]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Initialize a 4x4 matrix to hold confusion matrices for each model
combined_conf_matrix = np.zeros((4, 4), dtype=int)

# Define the size for the grid (2x2 confusion matrices for each model)
matrix_size = 2 # Since we're working with binary classification

# Add confusion matrices of each model to the 4x4 grid
for idx, (name, matrix) in enumerate(conf_matrices.items()):
    row_start = (idx // 2) * matrix_size # Determine the starting row
    col_start = (idx % 2) * matrix_size # Determine the starting column
    combined_conf_matrix[row_start:row_start+matrix_size, col_start:col_start+matrix_size] = matrix


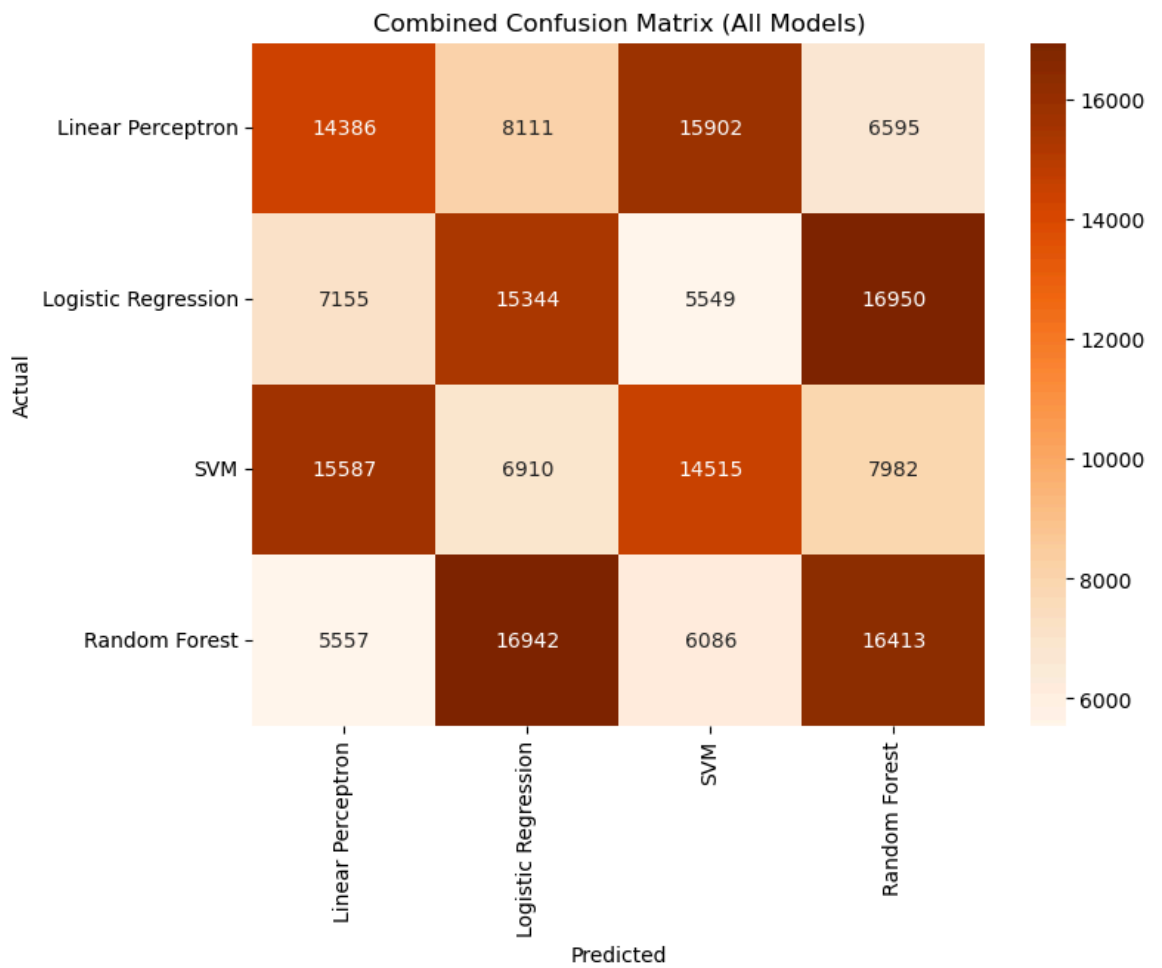
# Plotting the combined confusion matrix as a 4x4 matrix
plt.figure(figsize=(8, 6))
sns.heatmap(combined_conf_matrix, annot=True, fmt='d', cmap='Oranges',
            xticklabels=["Linear Perceptron", "Logistic Regression", "Support Vector Machine", "Random Forest"],
            yticklabels=["Linear Perceptron", "Logistic Regression", "Support Vector Machine", "Random Forest"])
plt.title("Combined Confusion Matrix (All Models)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Normalize the confusion matrices and combine them
normalized_matrices = [matrix / matrix.sum() for matrix in conf_matrices.values()]
combined_normalized = np.zeros((4, 4))

# Add normalized matrices to the combined normalized matrix
for idx, matrix in enumerate(normalized_matrices):
    row_start = (idx // 2) * matrix_size
    col_start = (idx % 2) * matrix_size
    combined_normalized[row_start:row_start+matrix_size, col_start:col_start+matrix_size] = matrix

# Plotting the combined normalized confusion matrix as a 4x4 matrix
plt.figure(figsize=(8, 6))
sns.heatmap(combined_normalized, annot=True, fmt='.2f', cmap='Purples',
            xticklabels=["Linear Perceptron", "Logistic Regression", "Support Vector Machine", "Random Forest"],
            yticklabels=["Linear Perceptron", "Logistic Regression", "Support Vector Machine", "Random Forest"])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```




```

In [17]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Define confusion matrices for each model
conf_matrices = {
    "Linear Perceptron": np.array([[14386, 8111], [15902, 6595]]),
    "Logistic Regression": np.array([[7155, 15344], [5549, 16950]]),
    "SVM": np.array([[15587, 6910], [14515, 7982]]),
    "Random Forest": np.array([[5557, 16942], [6086, 16413]])
}

model_names = list(conf_matrices.keys())
matrix_size = 2
grid_size = 4

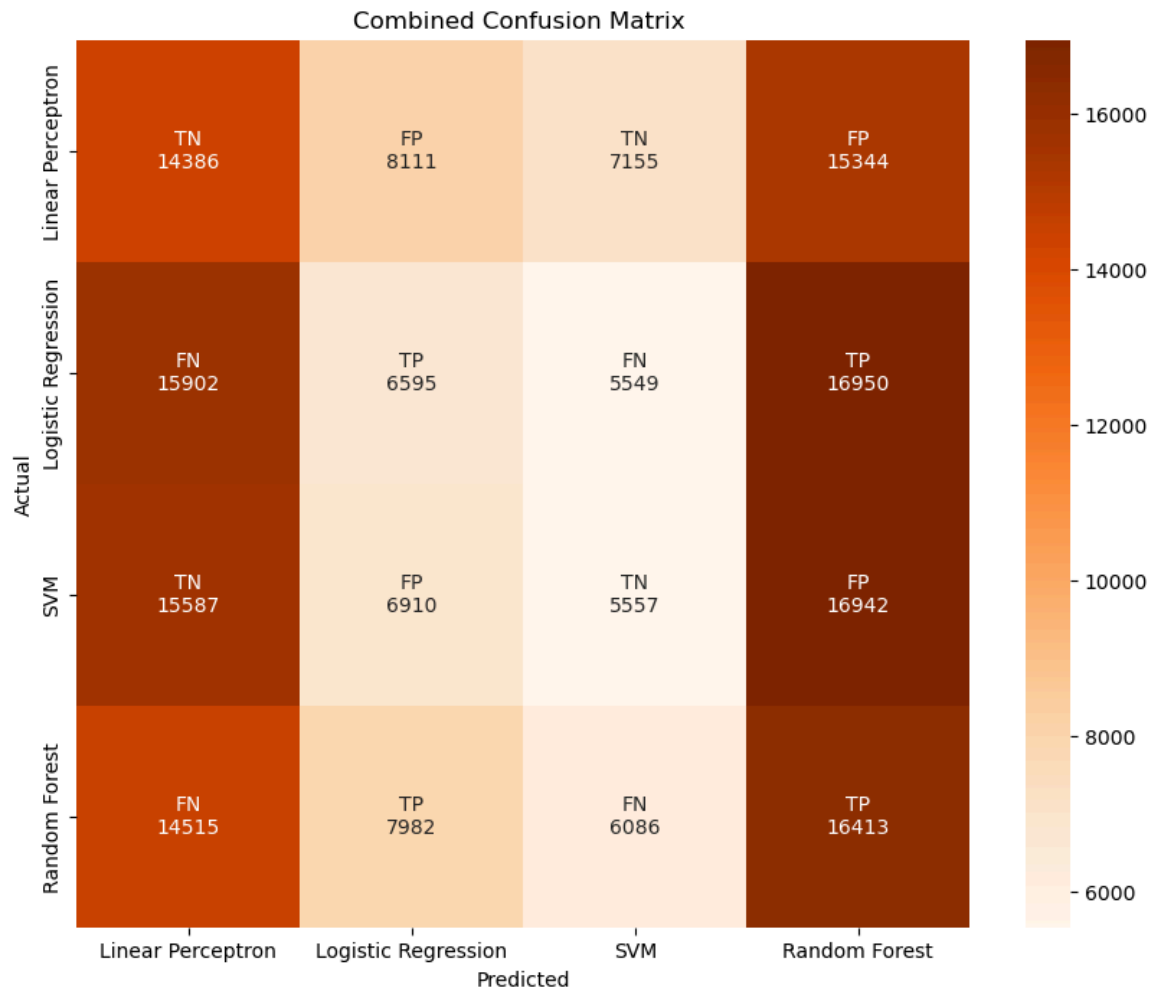
# Combined matrix and annotation text
combined_conf_matrix = np.zeros((grid_size, grid_size), dtype=int)
annotations = ["" for _ in range(grid_size)] for _ in range(grid_size)

# Fill combined matrix and annotation labels
for idx, (model, matrix) in enumerate(conf_matrices.items()):
    row_start = (idx // 2) * matrix_size
    col_start = (idx % 2) * matrix_size
    combined_conf_matrix[row_start:row_start+matrix_size, col_start:col_start+matrix_size] = matrix

    # Annotating each cell with its meaning
    annotations[row_start][col_start] = f"TN\n{matrix[0][0]}"
    annotations[row_start][col_start+1] = f"FP\n{matrix[0][1]}"
    annotations[row_start+1][col_start] = f"FN\n{matrix[1][0]}"
    annotations[row_start+1][col_start+1] = f"TP\n{matrix[1][1]}"

# Plotting the labeled combined confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(combined_conf_matrix, annot=annotations, fmt='', cmap='Oranges',
            xticklabels=model_names, yticklabels=model_names)
plt.title("Combined Confusion Matrix ")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```




```
In [18]: plt.figure(figsize=(8, 6))

colors = {
    "Perceptron": "red",
    "Logistic Regression": "blue",
    "SVM": "green",
    "Random Forest": "purple"
}

for name, model in models.items():
    if name in ["Perceptron", "SVM"]:
        y_scores = model.decision_function(X_val_tfidf) # Use decision function
    else:
        y_scores = model.predict_proba(X_val_tfidf)[:, 1] # Use predict_proba

    fpr, tpr, _ = roc_curve(y_val, y_scores)
    auc_score = auc(fpr, tpr)
    plt.plot(fpr, tpr, color=colors[name], label=f"{name} (AUC = {auc_score})")

plt.plot([0, 1], [0, 1], 'k--') # Diagonal line
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for All Models")
plt.legend()
plt.show()
```




```
In [19]: from sklearn.metrics import roc_curve, auc

roc_aucs = {}

for name, model in models.items():
    if name in ["Perceptron", "SVM"]:
        y_scores = model.decision_function(X_val_tfidf) # Use decision function
    else:
        y_scores = model.predict_proba(X_val_tfidf)[:, 1] # Use predict_proba

    fpr, tpr, _ = roc_curve(y_val, y_scores)
    roc_aucs[name] = auc(fpr, tpr)

# ===== 4 ROC AUC BAR PLOT =====
plt.figure(figsize=(6, 4))
sns.barplot(x=list(roc_aucs.keys()), y=list(roc_aucs.values()), palette="magma")
plt.title("ROC AUC Scores for All Models")
plt.ylabel("AUC Score")
plt.ylim(0.5, 1.0) # AUC ranges from 0.5 (random) to 1.0 (perfect)
plt.show()
```



```
In [20]: pipeline_file = open("text_sentiment.pkl", "wb")
joblib.dump(pipe_lr, pipeline_file)
pipeline_file.close()
```

```
In [21]: # Choose a model from previously trained models
model = models["Logistic Regression"] # or "SVM", "Random Forest", "K

# Define test sentences
test_sentences = [
    "My oldest son is making a trip with my mother... I'm missing him",
    "Heading off to bed... Goodnight peeps",
    "Comfortablity has won out",
    "@justineville ...yeahhh. ) i'm 39 tweets from 1,600!",
    "@buggin469 I'm jealous... Kinda miss you guys",
    "blaaah. I don't feel good again",
    "I can't believe I didn't get any sleep last night – not even a s
]

# Transform and predict
test_sentences_tfidf = vectorizer.transform(test_sentences)
predicted_sents = model.predict(test_sentences_tfidf)

# Display results
for i, (sentence, sentiment) in enumerate(zip(test_sentences, predicted_sents)):
    print(f"Sentence {i+1}: {sentence}")
    print(f"Predicted Sentiment: {'Positive' if sentiment == 1 else 'Negative'}
```

Sentence 1: My oldest son is making a trip with my mother... I'm missing him
Predicted Sentiment: Negative

Sentence 2: Heading off to bed... Goodnight peeps
Predicted Sentiment: Positive

Sentence 3: Comfortablity has won out
Predicted Sentiment: Positive

Sentence 4: @justineville ...yeahhh.) i'm 39 tweets from 1,600!
Predicted Sentiment: Positive

Sentence 5: @buggin469 I'm jealous... Kinda miss you guys
Predicted Sentiment: Negative

Sentence 6: blaaah. I don't feel good again
Predicted Sentiment: Negative

Sentence 7: I can't believe I didn't get any sleep last night – not even a single hour. Coffee isn't helping, and I can't focus on anything today. It feels like this day just doesn't want to cooperate. 😞
#Exhausted #RunningOnEmpty
Predicted Sentiment: Negative

In [22]: !pip install streamlit

```
python3.9/site-packages (from pandas<3,>=1.4.0->streamlit) (2024.1)
Requirement already satisfied: certifi>=2017.4.17 in ./opt/anaconda
3/lib/python3.9/site-packages (from requests<3,>=2.27->streamlit)
(2024.8.30)
Requirement already satisfied: charset-normalizer<4,>=2 in ./opt/ana
conda3/lib/python3.9/site-packages (from requests<3,>=2.27->stream
lit) (2.0.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in ./opt/anaconda
3/lib/python3.9/site-packages (from requests<3,>=2.27->streamlit)
(1.26.11)
Requirement already satisfied: idna<4,>=2.5 in ./opt/anaconda3/lib/
python3.9/site-packages (from requests<3,>=2.27->streamlit) (3.3)
Collecting smmap<6,>=3.0.1
  Using cached smmap-5.0.2-py3-none-any.whl (24 kB)
Requirement already satisfied: MarkupSafe>=0.23 in ./opt/anaconda3/
lib/python3.9/site-packages (from jinja2->altair<6,>=4.0->streamli
t) (2.0.1)
Requirement already satisfied: attrs>=17.4.0 in ./opt/anaconda3/li
b/python3.9/site-packages (from jsonschema>=3.0->altair<6,>=4.0->st
reamlit) (21.4.0)
```

In [23]: **from** sklearn.metrics **import** precision_score, f1_score

```
# Dictionaries to store precision and f1 scores
precisions = {}
f1_scores = {}

# Loop through models and predictions
for name, model in models.items():
    y_pred = model.predict(X_val_tfidf)
    precisions[name] = precision_score(y_val, y_pred)
    f1_scores[name] = f1_score(y_val, y_pred)

# Display as DataFrame
metrics_df = pd.DataFrame({
    'Precision': precisions,
    'F1 Score': f1_scores
})

print(metrics_df)
```

	Precision	F1 Score
Perceptron	0.654189	0.667798
Logistic Regression	0.719898	0.736252
SVM	0.710297	0.731031
Random Forest	0.672802	0.700004

```
In [24]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score

# Define model predictions (use the actual predictions of your models)
models = {
    'Perceptron': perceptron.predict(X_val_tfidf),
    'Logistic Regression': pipe_lr.predict(X_val),
    'SVM': pipe_svm.predict(X_val),
    'Random Forest': pipe_rf.predict(X_val)
}

model_names = list(models.keys())
n = len(model_names)

# Initialize the agreement matrix
agreement_matrix = np.zeros((n, n))

# Calculate the accuracy for each model and the agreement between the
for i in range(n):
    for j in range(n):
        if i == j:
            # Diagonal: accuracy of each model
            agreement_matrix[i, j] = accuracy_score(y_val, models[model_names[i]])
        else:
            # Off-diagonal: agreement between two models (proportion of
            match = np.sum(models[model_names[i]] == models[model_names[j]])
            agreement_matrix[i, j] = match / len(y_val)

# Normalize the matrix to ensure values between 0 and 1
agreement_matrix = np.clip(agreement_matrix, 0, 1)

# Plotting the normalized agreement matrix
plt.figure(figsize=(8, 6))
sns.heatmap(agreement_matrix, annot=True, fmt='.2f', cmap='Blues',
            xticklabels=model_names, yticklabels=model_names)
plt.title("Normalized Agreement Matrix (Accuracy and Agreement between")
plt.xlabel("Compared To")
plt.ylabel("Model")
plt.show()
```

Normalized Agreement Matrix (Accuracy and Agreement between Models)

